A NOTE FROM SHASHANK · APRIL 2026

What's next for Hatchin

— and where I need your help to get the order right.

Hi there,

Thanks for being on Hatchin. v2.0 just shipped and I'm sitting down to plan the next 6–9 months of the product. I have a roadmap. **It's plausible. It might also be wrong.** What I'm missing is your voice — the people actually using this every day.

This letter is short. Three things:

- What I'm thinking about building, in plain English
- The bigger swings I'm considering for later
- Three quick questions where your answer changes what I build first

If you read nothing else, jump to the black box on the last page and reply. Even one sentence helps.

— Shashank
Building Hatchin

Where Hatchin is right now

v2.0 ("Hatches That Deliver") shipped end of March. Your Hatch team can now produce coordinated deliverables across roles — content sprints, launch packages, research bundles. Before that, v1.3 made autonomy visible (the right sidebar), v1.2 brought billing and smarter LLM routing, v1.1 gave Hatches the ability to actually work in the background, not just chat.

That's the foundation. Now the question is what to build on top of it.

Seven things I'm considering, in plain English

Not all of these will ship. Not in this order. That's where you come in. Each one is a few weeks of work, and I want to know which actually solves a real problem for you.

1 Refinements that don't quietly make things worse

Right now, when you click "refine" on a Hatch's deliverable, the new version *could* be worse than the old one and you might not notice. I want to fix that — score every refinement against a fixed quality bar, and if the new version is worse, automatically keep the previous one and tell you why.

2 Hatches that look things up live, instead of guessing

Today your SEO Hatch (Robin), Growth Hatch (Kai), and Brand Hatch (Cass) work from training data — which is months out of date. I want them to actually search the live web mid-conversation, with real citations, before answering. Pro tier feature.

3 Upload PDFs, Excel, PowerPoint — Hatches actually read them

If you upload a scanned PDF today, Hatchin silently fails to read it. If you upload a spreadsheet with tables, the tables get mushed into garbage text. I want to fix all of that — proper OCR for scans, structured table reading, and forms a Hatch can fill out (NDAs, vendor questionnaires, expense reports — populated from your project brain).

4 Drop in a brief, get a full plan

Paste a project spec into Maya. She breaks it into a hierarchical task tree — who does what, what depends on what, what gets delivered. Hatches then work through it autonomously, with proper handoffs. It's the difference between "AI team that chats" and "AI team that actually plans and runs a project."

5 Marketing Hatches become real specialists, not generalists

Robin (SEO), Wren (Copy), Kai (Growth), Drew (Email), Pixel (Social), Cass (Brand) — today they do passable work. I want to give each of them the deep frameworks an actual senior practitioner uses. So Robin doesn't just say "improve your SEO," she walks you through a real schema markup audit. Wren doesn't just write copy — she writes copy that follows real conversion principles.

6 Real files, not just text descriptions

Today Cass gives you a markdown outline of a pitch deck. You then rebuild it in Keynote. With this: she gives you an actual **.pptx** file with proper slide design. Blake gives you a working **.xlsx** financial model, not a markdown table. Cleo gives you an interactive HTML prototype, not a description of one. Plus a brand layer so it all looks like *your* brand, not generic AI output.

7 Use your Hatch team from outside Hatchin

Right now your configured Hatches only work inside Hatchin's chat UI. I want to let your CMS call Robin on a schedule. Let your Slack bot route questions to Maya. Make every project an API endpoint other apps can call. Schedule Hatches on cron — "every Monday morning, ask Robin to audit our top 10 SEO pages and post results to #marketing."

Two bigger swings I'm thinking about for later

These are bigger and further out. I'm including them because if either of them is what you actually want, that changes everything else.

BIGGER SWING 1

Hatches that remember

Today, Hatches start each conversation cold. They don't remember the brand decision you locked in three weeks ago, the freemium pricing question Blake pushed back on last month, or the design direction Cleo championed in March. I want to fix this — give every Hatch persistent memory of your project, your decisions, your hard nos. So when you come back after a week, Maya opens with *"Cass is still concerned about the brand color drift you flagged. Blake is waiting on your call about freemium."* Instead of *"What were we working on?"*

BIGGER SWING 2

Hatches that actually do things, not just describe them

The current version of Hatchin tells you *what* to do. Robin writes the SEO audit; you fix the meta tags. Drew writes the cold email; you send it. I want Hatches to actually **do** the work — open Webflow, edit the meta tags, send the email, post to LinkedIn, file the Linear ticket. With your one-click approval each time. Mac-only at first. This is the biggest jump on the table — it changes Hatchin from "AI team that writes" to "AI team that ships." Probably 12 months out, but planting the flag now.

The roadmap is plausible.
Whether it's *actually* what you need — only you can tell me.

Three questions. Reply with whatever fits.

1. **Out of the seven things on the previous page, which TOP THREE would you want first?**
   Just numbers (e.g. "3, 6, 1") is fine. If something's missing entirely, name that instead.

2. **When you come back to Hatchin after a few days, does the Hatch remember enough — or do you find yourself re-explaining things?**
   Yes / no, plus one example if you have one. (This tells me how badly we need the "Hatches that remember" swing.)

3. **If a Hatch could literally do anything for you — even things that sound impossible — what would you most want?**
   Don't filter. (This tells me whether the "Hatches that do" swing is the right long bet, or whether I'm chasing the wrong horizon entirely.)

How to reply: just hit reply to this email, or message me directly on Hatchin (Maya can pass it along). One paragraph is plenty. I read every single one.

Why your input matters this time: Claude Code lets me ship features in 1-2 week cycles now. Whatever you tell me actually shows up in the product fast. This isn't a "we'll consider it" survey — it's an "I'm about to start building, please redirect me if I'm wrong" letter.

Thanks for using Hatchin. The fact that you've stuck with v2.0 means I owe you a v2.1 that's actually worth your time. Help me make sure it is.

— Shashank